7/Oct/24 (U0)

https://youtu.be/gMOAud7hZg4?list=PLSrTvUm384I9PV10koj_cqit9OfbJXEkq

Stanford, ML Sys

Video 67 — flash Attention (F.A)

Fast and Memory Efficient Exact Attention with IO Awareness

Why this?

3 main reasons for Modelling Longer Sequences

① NLP :- longer context is required to understand books, plays, ... etc

② CV :- higher resolution → better insights

③ Time series, Audio, Video, Medical Images → data understanding

↓

these are sequences of Million steps

Challenge:-

Scaling Transformers to longer Sequences :-

* context length of GPT 3 → 2048 (Seq)

F.A → helps to train things faster and with longer content.

① Tiling
② Recomputation
— 2 main things in F.A to reduce GPU memory IOs.

Speed → 3 times faster

Memory Efficient → 10 - 20 x for exact attention

Seq length → upto 16k (Context)

# Table of Content



## ① Background

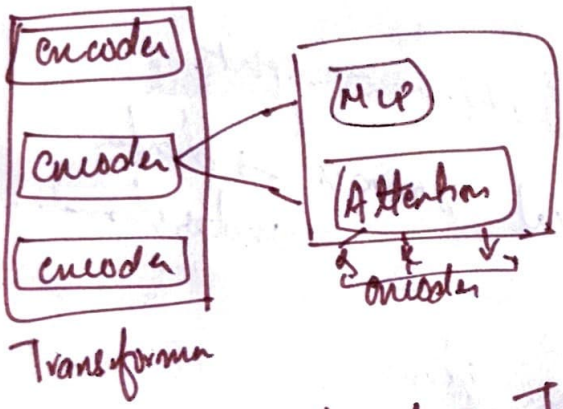


Q = Query
k = Key
v = Value
N = Seq length (2k, 4k, 16k)
d = head dimension (64, 128 ...)

→ Attention is heart of Transformer
→ Attention is bottlenecked by Memory Read/writes

$$O = \text{Dropout}(\text{Softmax}(\text{mask}[QK^T]))V$$

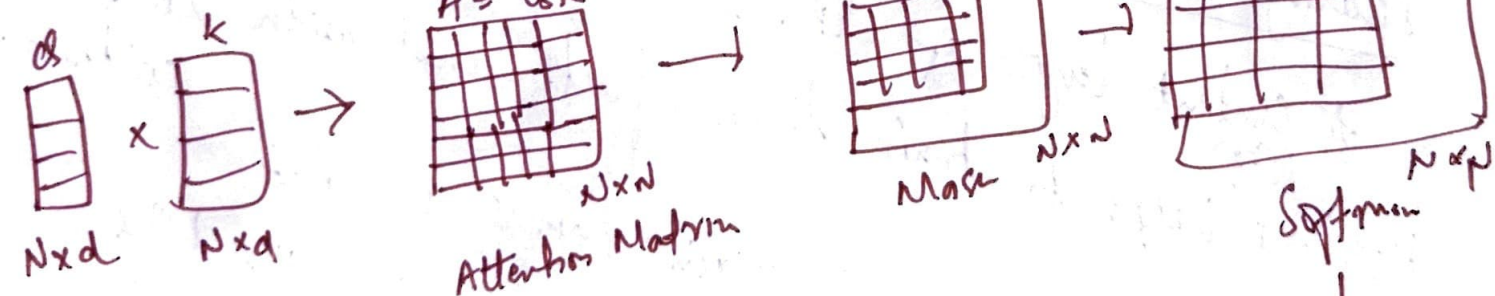


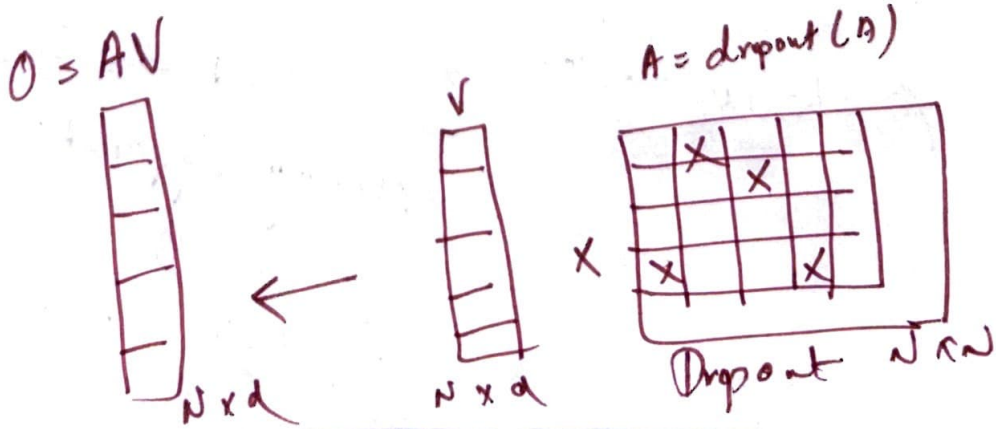

→ The problem is, lot of time is spent on Reading and writing in GPU instead of Calculations.

→ Attention in GPT 2 : Masking, softmax and Dropout taken lot of time than matrix multiplication.

→ hardware :- there is memory hierarchy in hardware

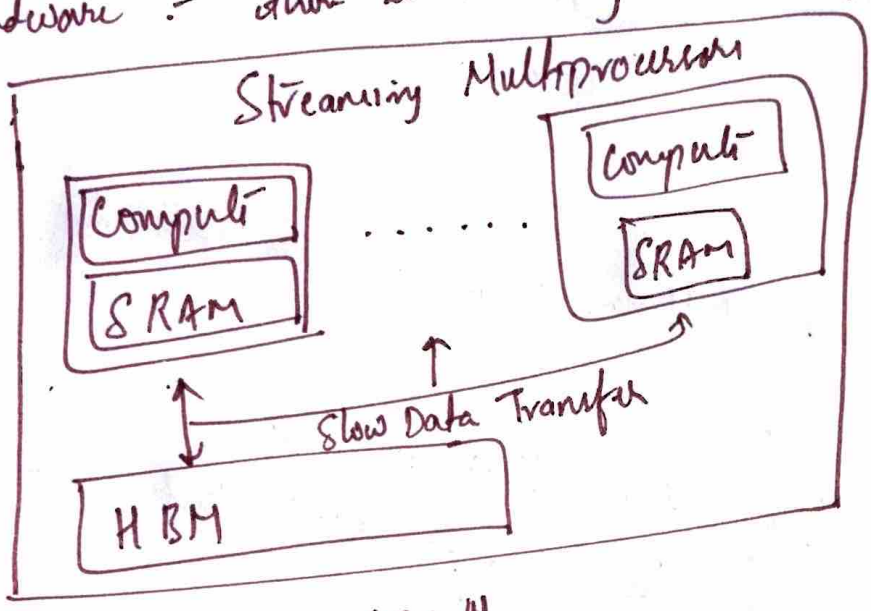


GPU #1

→ SRAM is more much smaller but much more faster

while any computation :-

inputs → HBM → SRAM ⇔ Compute → HBM

(Softmax, Matmul ...etc)

→ data moves back and forth between HBM and ~~GPU~~ SRAM

## ② Method : F.A

Challenges :- How to reduce HBM Read/Write : compute by blocks :-

① compute softmax reduction without access to full input

② backward without the large attention matrix from forward.

Techniques used to address these issues are:-

① Tiling : Restructure algorithm to load block by block from HBM to SRAM to compute attention

② Recomputation : Dont store attention matrix from forward, recompute it in the backward

Tiling :-

decomposing large softmax into smaller ones by scaling

Steps :-

① load inputs by block from HBM to SRAM
② On chip, compute attention outputs wrt that block
③ Update output in HBM by scaling

Recompute :-

by storing softmax normalization factors from fwd (size N) quickly recompute attention is load from inputs in SRAM